

Tomato disease classifier

This notebook is about building a tomato disease classifier that can be used to identify the possible diseases of your tomato using an image

```
# # downloaded the dataset:
# !curl -L -o /content/tomato-disease-multiple-sources.zip
https://www.kaggle.com/api/v1/datasets/download/cookiefinder/tomato-
disease-multiple-sources

# # Unzip the file
# !unzip /content/tomato-disease-multiple-sources.zip -d
/content/tomato-disease-dataset

# # Remove the zip file
# !rm /content/tomato-disease-multiple-sources.zip

# import os

# # Change the working directory
# os.chdir('/content/tomato-disease-dataset')

# # Verify the change
# print(os.getcwd())

# importing required libraries for the job

import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import transforms, models
from torch.utils.data import DataLoader, random_split
from torchvision.datasets import ImageFolder
import torchvision.transforms as T

import matplotlib.pyplot as plt
import numpy as np

# Check if any GPU is available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# If a CUDA device is available, use the first available GPU
if torch.cuda.is_available():
    device = torch.device("cuda:0") # Using the first available GPU
    at index 0
```

```

else:
    device = torch.device("cpu")

print(f"Using device: {device}")

Using device: cuda:0

global IMAGE_SIZE, BATCH_SIZE, CHANNELS, NUM_CLASSES, EPOCH
IMAGE_SIZE = 256
BATCH_SIZE = 16
CHANNELS = 3
EPOCH = 80

#defining the path of the dataset
data_path = '/kaggle/input/tomato-disease-multiple-sources/train'
#/content/tomato-disease-dataset/train

transform = transforms.Compose([
    transforms.Resize((IMAGE_SIZE, IMAGE_SIZE)), # Making sure that
the size is fixed for all
    transforms.ToTensor(), # Convert to tensor
])

#Load dataset from directory
dataset = ImageFolder(root=data_path, transform=transform)

# Create DataLoader with batching
data_loader = DataLoader(dataset, batch_size= BATCH_SIZE,
    shuffle=True,
                                num_workers=4,           # Number of
workers for data loading
                                pin_memory=True,         # Use pinned
                                memory for GPU transfer
                                prefetch_factor=2,       # Number of
                                batches to prefetch per workers
                                persistent_workers=True   # Keep workers
                                alive between epochs
                                )

CLASSES = dataset.classes
NUM_CLASSES = len(CLASSES)
print(NUM_CLASSES)

# View class names (subfolder names)
print("Class Names:")
for class_name in dataset.classes:
    print(class_name)

# View class-to-index mapping

```

```

print("\nClass-to-Index Mapping:")
for class_name, index in dataset.class_to_idx.items():
    print(f"{index}: {class_name}")

# View total number of samples
print("\nNumber of images:", len(dataset))

# View a single sample (image and label)
img1, label1 = dataset[0] # Access the first sample
print("Image shape:", img1.shape) # Shape of the image tensor
print("Label:", label1) # Integer label corresponding to the class

```

```

11
Class Names:
Bacterial_spot
Early_blight
Late_blight
Leaf_Mold
Septoria_leaf_spot
Spider_mites Two-spotted_spider_mite
Target_Spot
Tomato_Yellow_Leaf_Curl_Virus
Tomato_mosaic_virus
healthy
powdery_mildew

```

```

Class-to-Index Mapping:
0: Bacterial_spot
1: Early_blight
2: Late_blight
3: Leaf_Mold
4: Septoria_leaf_spot
5: Spider_mites Two-spotted_spider_mite
6: Target_Spot
7: Tomato_Yellow_Leaf_Curl_Virus
8: Tomato_mosaic_virus
9: healthy
10: powdery_mildew

```

```

Number of images: 25851
Image shape: torch.Size([3, 256, 256])
Label: 0

```

```

num = len(data_loader)
print(num)

```

```

1616

```

```

print(num*32) #Matches the number of batches we have created and the
number of images we have got

```

51712

```
for images, labels in data_loader:
    print(images.shape)
    print(labels.numpy())
    break

torch.Size([16, 3, 256, 256])
[ 9  0  7  9  9  8  4  5  7 10  0  6  7  9  0  2]

#for a singular image
for images, labels in data_loader:
    fig, axes = plt.subplots(3, 4, figsize=(7, 6))

    # Plot the first 12 images
    for i in range(12):
        row, col = divmod(i, 4)
        ax = axes[row, col]
        ax.imshow(images[i].permute(1, 2, 0)) # Change from C x H x W
to H x W x C)
        ax.set_title(f"Label: class {labels[i].item()}")
        ax.axis('off') # Turn off axes

plt.tight_layout()
plt.show()
break
```

Label: class 1



Label: class 9



Label: class 3



Label: class 6




Label: class 9



Label: class 6



Label: class 10



Label: class 3



Label: class 2



Label: class 5



Label: class 4



Label: class 6



```
#defining the path of the valid and test dataset
valid_data_path =
'/kaggle/input/tomato-disease-multiple-sources/valid'
#'/content/tomato-disease-dataset/valid'

# Define the transformation
valid_transform = transforms.Compose([
    transforms.Resize((IMAGE_SIZE, IMAGE_SIZE)), # Making sure that
the size is fixed for all
    transforms.ToTensor(), # Convert to tensor
])

# Load dataset from directory
valid_dataset = ImageFolder(root=valid_data_path,
transform=valid_transform)

# Define the split sizes
valid_size = int(0.5 * len(valid_dataset)) # 50% for validation
```

```

test_size = len(valid_dataset) - valid_size # Remaining for test

# Split the dataset
valid_subset, test_subset = random_split(valid_dataset, [valid_size,
test_size])

# Create DataLoaders for validation and test sets
# Set a seed for reproducibility
torch.manual_seed(42)
valid_loader = DataLoader(valid_subset, batch_size= BATCH_SIZE,
shuffle=True,
                                num_workers=4,           # Number of
workers for data loading
                                pin_memory=True,         # Use pinned
memory for GPU transfer
                                prefetch_factor=2,       # Number of
batches to prefetch per workers
                                persistent_workers=True   # Keep workers
alive between epochs
                                )

# Set a seed for reproducibility
torch.manual_seed(42)
test_loader = DataLoader(test_subset, batch_size= BATCH_SIZE,
shuffle=True,
                                num_workers=4,           # Number of
workers for data loading
                                pin_memory=True,         # Use pinned
memory for GPU transfer
                                prefetch_factor=2,       # Number of
batches to prefetch per workers
                                persistent_workers=True   # Keep workers
alive between epochs
                                )

# Check the size of SETS
print(f"Training set size: {len(data_loader)}")
print(f"Validation set size: {len(valid_loader)}")
print(f"Test set size: {len(test_loader)}")

Training set size: 1616
Validation set size: 209
Test set size: 209

#Checking whether the datasets are fine
for images, labels in valid_loader:

```

```

print(images[0].shape)
print(labels[0].numpy())
plt.imshow(images[0].permute(1, 2, 0)) # Change from C x H x W to
H x W x C)
plt.title(f"Valid dataset \n Label: class {labels[0]}")
plt.axis('off')
plt.show()
break

torch.Size([3, 256, 256])
5

```

Valid dataset
Label: class 5



```

for images, labels in test_loader:
    print(images[0].shape)
    print(labels[0].numpy())
    plt.imshow(images[0].permute(1, 2, 0)) # Change from C x H x W to
H x W x C)
    plt.title(f"Test Dataset \nLabel: class {labels[0]}")
    plt.axis('off')
    plt.show()
    break

torch.Size([3, 256, 256])
7

```


Test Dataset
Label: class 7



```
# For scaling and resizing incase external images are used
class ScaleAndResizeLayer(nn.Module):
    def __init__(self, size=(IMAGE_SIZE, IMAGE_SIZE),
scale=1.0/255.0):
        super(ScaleAndResizeLayer, self).__init__()
        self.resize = T.Resize(size) # Resize to the target size
        self.scale = scale # Scale factor

    def forward(self, x):
        # Resize and scale
        x = self.resize(x)
        x = x * self.scale # Scale the pixel values
        return x

#Augmenting the input the data for better modelling
class AugmentationLayer(nn.Module):
    def __init__(self, image_size=(IMAGE_SIZE, IMAGE_SIZE)):
        super(AugmentationLayer, self).__init__()
        self.augmentations = T.Compose([
            T.RandomHorizontalFlip(p=0.5), # 50% chance of
horizontal flip
            T.RandomRotation(degrees=30), # Random rotation
within ±30 degrees
            T.Resize(image_size), # Resize to the
```



```

target size
T.RandomCrop(image_size),          # Optional:
Random crop
T.ColorJitter(brightness=0.2, contrast=0.2,
saturation=0.2, hue=0.1), # Random color jitter
])

def forward(self, x):
    # Apply augmentations
    return self.augmentations(x)

#Building the CNN -based model now

class MyModel(nn.Module):
    def __init__(self):
        super(MyModel, self).__init__()
        self.model = nn.Sequential(
            # Preprocessing and augmentation layers
            ScaleAndResizeLayer(),          #
Scaling and Resizing layer
            AugmentationLayer(),            #
Augmentation layer

            # First convolution block
            nn.Conv2d(3, 32, kernel_size=(3, 3), stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

            # Additional convolution blocks
            nn.Conv2d(32, 64, kernel_size=(3, 3), stride=1,
padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(64, 64, kernel_size=(3, 3), stride=1,
padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(64, 64, kernel_size=(3, 3), stride=1,
padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

```

```

        nn.Conv2d(64, 64, kernel_size=(3, 3), stride=1,
padding=1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2),

        # Flatten layer
        nn.Flatten(),

        # Fully connected (dense) layer
        nn.Linear(64 * 4 * 4, 64), # Adjust input features based
on the final dimensions
        nn.ReLU(),

        # Final classification layer
        nn.Linear(64, NUM_CLASSES), # Final layer with number of
neurons equal to number of classes
        nn.Softmax(dim=1) # Softmax for multi-class
classification
    )

    def forward(self, x):
        assert x.shape[1:] == (3, 256, 256), f"Expected (3, 256, 256),
got {x.shape[1:]}" # Check input shape
        return self.model(x)

```

`!pip install torchsummary` # Install the torchsummary package

#Printing a summary of the built model

```
from torchsummary import summary
```

```
model = MyModel()
```

```
model.to(device)
```

```
summary(model, input_size=(3, 256, 256)) # (Channels, Height, Width)
```

Requirement already satisfied: torchsummary in
/usr/local/lib/python3.10/dist-packages (1.5.1)

Layer (type)	Output Shape	Param #
Resize-1	[-1, 3, 256, 256]	0
ScaleAndResizeLayer-2	[-1, 3, 256, 256]	0
AugmentationLayer-3	[-1, 3, 256, 256]	0
Conv2d-4	[-1, 32, 256, 256]	896
ReLU-5	[-1, 32, 256, 256]	0
MaxPool2d-6	[-1, 32, 128, 128]	0
Conv2d-7	[-1, 64, 128, 128]	18,496
ReLU-8	[-1, 64, 128, 128]	0

MaxPool2d-9	[-1, 64, 64, 64]	0
Conv2d-10	[-1, 64, 64, 64]	36,928
ReLU-11	[-1, 64, 64, 64]	0
MaxPool2d-12	[-1, 64, 32, 32]	0
Conv2d-13	[-1, 64, 32, 32]	36,928
ReLU-14	[-1, 64, 32, 32]	0
MaxPool2d-15	[-1, 64, 16, 16]	0
Conv2d-16	[-1, 64, 16, 16]	36,928
ReLU-17	[-1, 64, 16, 16]	0
MaxPool2d-18	[-1, 64, 8, 8]	0
Conv2d-19	[-1, 64, 8, 8]	36,928
ReLU-20	[-1, 64, 8, 8]	0
MaxPool2d-21	[-1, 64, 4, 4]	0
Flatten-22	[-1, 1024]	0
Linear-23	[-1, 64]	65,600
ReLU-24	[-1, 64]	0
Linear-25	[-1, 11]	715
Softmax-26	[-1, 11]	0

=====
Total params: 233,419

Trainable params: 233,419

Non-trainable params: 0

Input size (MB): 0.75

Forward/backward pass size (MB): 64.49

Params size (MB): 0.89

Estimated Total Size (MB): 66.13

'''

This model in it self when trained results in very low accuracies. Hence, I thought of initiating with pre trained weights to have a better chgen of having a good accuracy

'''

```
# def train_model(model, train_loader, val_loader, criterion,
optimizer, num_epochs, device):
```

```
#
```

```
# """
#     Trains and validates a PyTorch model, similar to model.fit in
Keras.
```

```
#     Args:
```

```
#         model: The PyTorch model to train.
```

```
#         train_loader: DataLoader for training data.
```

```
#         val_loader: DataLoader for validation data.
```

```
#         criterion: Loss function.
```

```
#         optimizer: Optimizer.
```

```
#         num_epochs: Number of epochs to train.
```

```
#         device: Device to run the model on ('cpu' or 'cuda').
```

```

# Returns:
#     Trained model.
# """
# for epoch in range(num_epochs):
#     # Training phase
#     model.train() # Set model to training mode
#     train_loss = 0.0
#     train_correct = 0
#     train_total = 0

#     for inputs, labels in train_loader:
#         inputs, labels = inputs.to(device), labels.to(device)

#         # Zero the parameter gradients
#         optimizer.zero_grad()

#         # Forward pass
#         outputs = model(inputs)

#         # Compute loss
#         loss = criterion(outputs, labels)

#         # Backward pass and optimization
#         loss.backward()
#         optimizer.step()

#         # Accumulate loss and metrics
#         train_loss += loss.item() * inputs.size(0)
#         train_correct += (outputs.argmax(1) ==
# labels).sum().item()
#         train_total += labels.size(0)

#     # Calculate average training loss and accuracy
#     avg_train_loss = train_loss / train_total
#     train_accuracy = train_correct / train_total

#     # Validation phase
#     model.eval() # Set model to evaluation mode
#     val_loss = 0.0
#     val_correct = 0
#     val_total = 0

#     with torch.no_grad():
#         for inputs, labels in val_loader:
#             inputs, labels = inputs.to(device),
# labels.to(device)

#             # Forward pass
#             outputs = model(inputs)

```

```

#             # Compute loss
#             loss = criterion(outputs, labels)

#             # Accumulate loss and metrics
#             val_loss += loss.item() * inputs.size(0)
#             val_correct += (outputs.argmax(1) ==
labels).sum().item()
#             val_total += labels.size(0)

#         # Calculate average validation loss and accuracy
#         avg_val_loss = val_loss / val_total
#         val_accuracy = val_correct / val_total

#         # Print progress for the current epoch
#         print(f"Epoch {epoch+1}/{num_epochs}")
#         print(f"  Train Loss: {avg_train_loss:.4f}, Train Accuracy:
{train_accuracy:.4f}")
#         print(f"  Val Loss: {avg_val_loss:.4f}, Val Accuracy:
{val_accuracy:.4f}")

#     return model

```

'\nThis model in it self when trained results in very low accuracies.\nHence, I thought of initiating with pre trained weights to have a better chgen of having a good accuracy\n\n'

```

from tqdm.notebook import tqdm
import gc
from datetime import datetime
import warnings
warnings.filterwarnings('ignore')

class EarlyStopping:
    def __init__(self, patience=7, min_delta=0,
restore_best_weights=True):
        self.patience = patience
        self.min_delta = min_delta
        self.restore_best_weights = restore_best_weights
        self.best_model = None
        self.best_loss = None
        self.counter = 0
        self.status = ""

    def __call__(self, val_loss, model):
        if self.best_loss is None:
            self.best_loss = val_loss
            self.best_model = model.state_dict().copy()
        elif val_loss > self.best_loss - self.min_delta:

```

```

        self.counter += 1
        self.status = f'EarlyStopping counter: {self.counter} out
of {self.patience}'
        if self.counter >= self.patience:
            self.status = f'EarlyStopping triggered after
{self.counter} epochs'
            return True
        else:
            self.best_loss = val_loss
            self.best_model = model.state_dict().copy()
            self.counter = 0
            self.status = 'Improved'
        return False

def train_model(model, train_loader, val_loader, criterion, optimizer,
num_epochs, device,
                scheduler=None, early_stopping_patience=5):
    """
    Trains and validates a PyTorch model with Colab-optimized
features.

    Args:
        model: The PyTorch model to train
        train_loader: DataLoader for training data
        val_loader: DataLoader for validation data
        criterion: Loss function
        optimizer: Optimizer
        num_epochs: Number of epochs to train
        device: Device to run the model on ('cpu' or 'cuda')
        scheduler: Optional learning rate scheduler
        early_stopping_patience: Number of epochs to wait before early
stopping

    Returns:
        Trained model and dictionary containing training history
    """
    # Initialize early stopping
    early_stopping = EarlyStopping(patience=early_stopping_patience)

    # Initialize history dictionary
    history = {
        'train_loss': [], 'val_loss': [],
        'train_acc': [], 'val_acc': []
    }

    # Get the start time
    start_time = datetime.now()

    print(f"Training started at {start_time.strftime('%H:%M:%S')}")
    print(f"Using device: {device}")

```



```

try:
    for epoch in range(num_epochs):
        # Training phase
        model.train()
        train_loss = 0.0
        train_correct = 0
        train_total = 0

        # Use tqdm for progress bar
        train_pbar = tqdm(train_loader, desc=f'Epoch
{epoch+1}/{num_epochs} [Train]')

        for inputs, labels in train_pbar:
            inputs, labels = inputs.to(device), labels.to(device)

            # Clear gradients
            optimizer.zero_grad()

            # Forward pass
            outputs = model(inputs)
            loss = criterion(outputs, labels)

            # Backward pass and optimization
            loss.backward()
            optimizer.step()

            # Accumulate metrics
            train_loss += loss.item() * inputs.size(0)
            train_correct += (outputs.argmax(1) ==
labels).sum().item()
            train_total += labels.size(0)

            # Update progress bar
            train_pbar.set_postfix({
                'loss': f'{loss.item():.4f}',
                'acc': f'{train_correct/train_total:.4f}'
            })

        # Validation phase
        model.eval()
        val_loss = 0.0
        val_correct = 0
        val_total = 0

        val_pbar = tqdm(val_loader, desc=f'Epoch
{epoch+1}/{num_epochs} [Valid]')

        with torch.no_grad():
            for inputs, labels in val_pbar:

```

```

        inputs, labels = inputs.to(device),
labels.to(device)

        outputs = model(inputs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * inputs.size(0)
        val_correct += (outputs.argmax(1) ==
labels).sum().item()
        val_total += labels.size(0)

        val_pbar.set_postfix({
            'loss': f'{loss.item():.4f}',
            'acc': f'{val_correct/val_total:.4f}'
        })

    # Calculate epoch metrics
    avg_train_loss = train_loss / train_total
    avg_val_loss = val_loss / val_total
    train_accuracy = train_correct / train_total
    val_accuracy = val_correct / val_total

    # Update history
    history['train_loss'].append(avg_train_loss)
    history['val_loss'].append(avg_val_loss)
    history['train_acc'].append(train_accuracy)
    history['val_acc'].append(val_accuracy)

    # Print epoch summary
    print(f"\nEpoch {epoch+1}/{num_epochs} Summary:")
    print(f"Train Loss: {avg_train_loss:.4f}, Train Acc:
{train_accuracy:.4f}")
    print(f"Val Loss: {avg_val_loss:.4f}, Val Acc:
{val_accuracy:.4f}")

    # Learning rate scheduler step
    if scheduler is not None:
        scheduler.step()
        print(f"Learning Rate: {optimizer.param_groups[0]
['lr']:.6f}")

    # Early stopping check
    if early_stopping(avg_val_loss, model):
        print("Early stopping triggered!")
        model.load_state_dict(early_stopping.best_model)
        break

    # Clear GPU memory
    torch.cuda.empty_cache()
    gc.collect()

```

```

except KeyboardInterrupt:
    print("\nTraining interrupted by user!")
    if early_stopping.best_model is not None:
        print("Loading best model weights...")
        model.load_state_dict(early_stopping.best_model)

    # Calculate training time
    training_time = datetime.now() - start_time
    print(f"\nTraining completed in {training_time}")

    return model, history

# Initialize model, criterion, and optimizer with improved settings
for Colab
def initialize_training(model, learning_rate=1e-4):
    """
    Initialize training components with Colab-optimized settings
    """
    # Move model to GPU if available
    device = torch.device('cuda' if torch.cuda.is_available() else
'cpu')
    model = model.to(device)

    # Use mixed precision for faster training
    scaler = torch.cuda.amp.GradScaler()

    # Loss function with label smoothing for better generalization
    criterion = nn.CrossEntropyLoss(label_smoothing=0.1)

    # AdamW optimizer with weight decay
    optimizer = torch.optim.AdamW(model.parameters(),
                                   lr=learning_rate,
                                   weight_decay=0.01)

    # Cosine annealing scheduler with warm restarts
    scheduler = torch.optim.lr_scheduler.CosineAnnealingWarmRestarts(
        optimizer,
        T_0=10, # Reset LR every 10 epochs
        T_mult=2, # Double the reset interval after each restart
        eta_min=1e-6 # Minimum learning rate
    )

    return device, criterion, optimizer, scheduler, scaler

class PretrainedClassifier(nn.Module):
    def __init__(self, num_classes, model_name='resnet50',
freeze_backbone=True):
        """
        Initialize model with pre-trained backbone

```

```

Args:
    num_classes: Number of classes to predict
    model_name: Name of the pre-trained model to use
    freeze_backbone: Whether to freeze the backbone layers
"""
super(PretrainedClassifier, self).__init__()

# Available pretrained models
self.available_models = {
    'resnet18': models.resnet18,
    'resnet50': models.resnet50,
    'efficientnet_b0': models.efficientnet_b0,
    'mobilenet_v3_large': models.mobilenet_v3_large,
    'convnext_small': models.convnext_small
}

if model_name not in self.available_models:
    raise ValueError(f"Model {model_name} not available.
Choose from: {list(self.available_models.keys())}")

# Load pretrained model
self.backbone = self.available_models[model_name]
(pretrained=True)

# Freeze backbone if specified
if freeze_backbone:
    for param in self.backbone.parameters():
        param.requires_grad = False

# Replace the final layer based on model type
if model_name.startswith('resnet'):
    in_features = self.backbone.fc.in_features
    self.backbone.fc = nn.Sequential(
        nn.Linear(in_features, 512),
        nn.BatchNorm1d(512),
        nn.ReLU(),
        nn.Dropout(0.3),
        nn.Linear(512, num_classes)
    )
elif model_name.startswith('efficientnet'):
    in_features = self.backbone.classifier[-1].in_features
    self.backbone.classifier = nn.Sequential(
        nn.Linear(in_features, 512),
        nn.BatchNorm1d(512),
        nn.ReLU(),
        nn.Dropout(0.3),
        nn.Linear(512, num_classes)
    )
elif model_name.startswith('mobilenet'):

```

```

        in_features = self.backbone.classifier[-1].in_features
        self.backbone.classifier = nn.Sequential(
            nn.Linear(in_features, 512),
            nn.BatchNorm1d(512),
            nn.ReLU(),
            nn.Dropout(0.3),
            nn.Linear(512, num_classes)
        )
    elif model_name.startswith('convnext'):
        in_features = self.backbone.classifier[-1].in_features
        self.backbone.classifier = nn.Sequential(
            nn.Linear(in_features, 512),
            nn.LayerNorm(512),
            nn.GELU(),
            nn.Dropout(0.3),
            nn.Linear(512, num_classes)
        )

    def forward(self, x):
        return self.backbone(x)

def initialize_training_pretrained(model, num_epochs, train_loader,
learning_rate=1e-4):
    """
    Initialize training components optimized for fine-tuning
    """
    device = torch.device("cuda" if torch.cuda.is_available() else
"cpu")
    model = model.to(device)

    # Different learning rates for pretrained and new layers
    params = [
        {'params': [p for n, p in model.named_parameters() if
'backbone' not in n or 'fc' in n or 'classifier' in n],
         'lr': learning_rate},
        {'params': [p for n, p in model.named_parameters() if
'backbone' in n and 'fc' not in n and 'classifier' not in n],
         'lr': learning_rate/10}
    ]

    # Optimizer with weight decay
    optimizer = torch.optim.AdamW(params, weight_decay=0.01)

    # Loss function with label smoothing
    criterion = nn.CrossEntropyLoss(label_smoothing=0.1)

    # Learning rate scheduler
    scheduler = torch.optim.lr_scheduler.OneCycleLR(
        optimizer,
        max_lr=[learning_rate, learning_rate/10],

```

```

        epochs=num_epochs,
        steps_per_epoch=len(train_loader),
        pct_start=0.3,
        anneal_strategy='cos'
    )

    # Gradient scaler for mixed precision training
    scaler = torch.cuda.amp.GradScaler()

    return device, criterion, optimizer, scheduler, scaler

# Creating a model with pretrained weights

model = PretrainedClassifier(
    num_classes= NUM_CLASSES,
    model_name='resnet50',
    freeze_backbone=True    # freeze pretrained weights initially
)

# Initialize training components
device, criterion, optimizer, scheduler, scaler =
initialize_training_pretrained(
    model,
    num_epochs=EPOCH,
    train_loader=data_loader,
    learning_rate=1e-4
)

# Train as before
trained_model, history = train_model(
    model=model,
    train_loader=data_loader,
    val_loader=valid_loader,
    criterion=criterion,
    optimizer=optimizer,
    num_epochs=EPOCH,
    device=device,
    scheduler=scheduler,
    early_stopping_patience=5
)

Downloading: "https://download.pytorch.org/models/resnet50-
0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-
0676ba61.pth
100%|██████████| 97.8M/97.8M [00:00<00:00, 223MB/s]

Training started at 03:22:11
Using device: cuda

```



```
{"model_id": "25cee0437c1a402ca0d300b7c26c4179", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9f740369012442f2b83dcd683f4b756c", "version_major": 2, "version_minor": 0}
```

Epoch 1/80 Summary:

Train Loss: 1.8778, Train Acc: 0.4621

Val Loss: 1.4052, Val Acc: 0.7316

Learning Rate: 0.000004

```
{"model_id": "c1927e10bd5c48e7893dd12e7ca40133", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a1f9e6589c6b438693a9c0a40cd2b341", "version_major": 2, "version_minor": 0}
```

Epoch 2/80 Summary:

Train Loss: 1.3769, Train Acc: 0.7012

Val Loss: 1.1670, Val Acc: 0.7944

Learning Rate: 0.000004

```
{"model_id": "8a636a48a10c441687929c470ebac982", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0fdb911b4754fff81f9841ab6e5f49f", "version_major": 2, "version_minor": 0}
```

Epoch 3/80 Summary:

Train Loss: 1.2182, Train Acc: 0.7473

Val Loss: 1.0679, Val Acc: 0.8181

Learning Rate: 0.000004

```
{"model_id": "fef27562bbe54321a9384283df1a9eae", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "27366917de7c471db4598a0f9b86135c", "version_major": 2, "version_minor": 0}
```

Epoch 4/80 Summary:

Train Loss: 1.1421, Train Acc: 0.7719

Val Loss: 1.0161, Val Acc: 0.8303

Learning Rate: 0.000004

```
{"model_id": "234b1c6f77d74e77af2ddc7dc914e947", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a380cf99236041d381ba085cab7af22d", "version_major": 2, "version_minor": 0}
```

Epoch 5/80 Summary:

Train Loss: 1.0902, Train Acc: 0.7907

Val Loss: 0.9650, Val Acc: 0.8534

Learning Rate: 0.000004

```
{"model_id": "f989e4808f9745d4bd577764f93c9029", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "4238798d500948e6a87c307257656bb3", "version_major": 2, "version_minor": 0}
```

Epoch 6/80 Summary:

Train Loss: 1.0541, Train Acc: 0.8029

Val Loss: 0.9376, Val Acc: 0.8606

Learning Rate: 0.000004

```
{"model_id": "5b195ee52d8a44a19830f802d0c5a238", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5c6d815c9a144edeab049029a7ba75db", "version_major": 2, "version_minor": 0}
```

Epoch 7/80 Summary:

Train Loss: 1.0258, Train Acc: 0.8164

Val Loss: 0.9186, Val Acc: 0.8698

Learning Rate: 0.000004

```
{"model_id": "949eb9ae41894eb08a639cc2c08b0605", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "1520aa8487374ec4b284291a8cd9c197", "version_major": 2, "version_minor": 0}
```

Epoch 8/80 Summary:

Train Loss: 1.0080, Train Acc: 0.8214

Val Loss: 0.9115, Val Acc: 0.8645

Learning Rate: 0.000004

```
{"model_id": "0c2a92b7a6ab477682d75e03382e490d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e2498e0abea04bf0bb8450d60d90a650", "version_major": 2, "version_minor": 0}
```

Epoch 9/80 Summary:

Train Loss: 0.9883, Train Acc: 0.8296

Val Loss: 0.8950, Val Acc: 0.8749
Learning Rate: 0.000004

```
{"model_id": "13e27e318bba463d9abae259b667b323", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "bd436358739b4d2c936b41caf70c9d2f", "version_major": 2, "version_minor": 0}
```

Epoch 10/80 Summary:
Train Loss: 0.9719, Train Acc: 0.8359
Val Loss: 0.8853, Val Acc: 0.8752
Learning Rate: 0.000004

```
{"model_id": "2affad513ba245db9d1e429586f7bda6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5166222c17184a09a7441c61b60fab4a", "version_major": 2, "version_minor": 0}
```

Epoch 11/80 Summary:
Train Loss: 0.9597, Train Acc: 0.8405
Val Loss: 0.8676, Val Acc: 0.8845
Learning Rate: 0.000004

```
{"model_id": "11d62a72acbf4df7aa34f8bc922866a8", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "7ba1363c2fe64dd98dcb90060c9d7486", "version_major": 2, "version_minor": 0}
```

Epoch 12/80 Summary:
Train Loss: 0.9488, Train Acc: 0.8485
Val Loss: 0.8627, Val Acc: 0.8839
Learning Rate: 0.000004

```
{"model_id": "0d33c207d26a4bd582661ea78d4b38d6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "93a891579a8044ea90de75649e6c8b45", "version_major": 2, "version_minor": 0}
```

Epoch 13/80 Summary:
Train Loss: 0.9386, Train Acc: 0.8534
Val Loss: 0.8524, Val Acc: 0.8941
Learning Rate: 0.000004

```
{"model_id": "9f266626022d4db8a342b0358a58b550", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6834344fd907450e868b9fec5c58b165", "version_major": 2, "version_minor": 0}
```

Epoch 14/80 Summary:

Train Loss: 0.9247, Train Acc: 0.8566

Val Loss: 0.8513, Val Acc: 0.8923

Learning Rate: 0.000004

```
{"model_id": "ee76c4847b3b431e9951cae71e1a7948", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "41d58d1aae7c418788f9b496c2fd1391", "version_major": 2, "version_minor": 0}
```

Epoch 15/80 Summary:

Train Loss: 0.9176, Train Acc: 0.8626

Val Loss: 0.8471, Val Acc: 0.8911

Learning Rate: 0.000004

```
{"model_id": "a56a454b0c8e4cce88c99d5640812bb7", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "ad7db578e5194e6292dfcd0150935a94", "version_major": 2, "version_minor": 0}
```

Epoch 16/80 Summary:

Train Loss: 0.9079, Train Acc: 0.8648

Val Loss: 0.8266, Val Acc: 0.9078

Learning Rate: 0.000004

```
{"model_id": "ebb78448a9024f92880e1992ff3136d3", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "1798b87da8b743b68f6c8d317487f840", "version_major": 2, "version_minor": 0}
```

Epoch 17/80 Summary:

Train Loss: 0.9006, Train Acc: 0.8696

Val Loss: 0.8227, Val Acc: 0.9045

Learning Rate: 0.000004

```
{"model_id": "5db04f172704436ba73be12f09ab4111", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "2fa82a6934ba4f92ba963579295d5bd6", "version_major": 2, "version_minor": 0}
```

Epoch 18/80 Summary:

Train Loss: 0.8942, Train Acc: 0.8711
Val Loss: 0.8184, Val Acc: 0.9057
Learning Rate: 0.000004

```
{"model_id": "86d4049ec7584ab4865132d7a1a1869c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a228801065c744579aa1e2f682fe9ac7", "version_major": 2, "version_minor": 0}
```

Epoch 19/80 Summary:
Train Loss: 0.8869, Train Acc: 0.8755
Val Loss: 0.8119, Val Acc: 0.9102
Learning Rate: 0.000004

```
{"model_id": "d218102c514c4b58abddff785eb897d3", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5a2b447ce0aa40be8c3e58e6b945757c", "version_major": 2, "version_minor": 0}
```

Epoch 20/80 Summary:
Train Loss: 0.8803, Train Acc: 0.8770
Val Loss: 0.8128, Val Acc: 0.9081
Learning Rate: 0.000004

```
{"model_id": "2403c13611054606bd9e3bc42a917879", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "4499a7f8ce48474bb579944373b05dde", "version_major": 2, "version_minor": 0}
```

Epoch 21/80 Summary:
Train Loss: 0.8763, Train Acc: 0.8807
Val Loss: 0.8016, Val Acc: 0.9192
Learning Rate: 0.000004

```
{"model_id": "a4d43e90de354552ad6d726520019150", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e61d7589cd9c46dcbe8f258c727f1791", "version_major": 2, "version_minor": 0}
```

Epoch 22/80 Summary:
Train Loss: 0.8699, Train Acc: 0.8843
Val Loss: 0.7988, Val Acc: 0.9174
Learning Rate: 0.000004

```
{"model_id": "c21d236987fa41158ed423bdd369518c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "54122c1e449341b2ad5b3812ab25f811", "version_major": 2, "version_minor": 0}
```

Epoch 23/80 Summary:

Train Loss: 0.8652, Train Acc: 0.8857

Val Loss: 0.8032, Val Acc: 0.9147

Learning Rate: 0.000004

```
{"model_id": "337a9f42bdcc4e8e849a185b3b584071", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "15306839d0de42738f4db0899ef9b07e", "version_major": 2, "version_minor": 0}
```

Epoch 24/80 Summary:

Train Loss: 0.8603, Train Acc: 0.8882

Val Loss: 0.7888, Val Acc: 0.9234

Learning Rate: 0.000004

```
{"model_id": "3044797ae2394a0d957c27002c94d8b3", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "ffc11afd3ca2436b90e7b19c5e9803ef", "version_major": 2, "version_minor": 0}
```

Epoch 25/80 Summary:

Train Loss: 0.8536, Train Acc: 0.8921

Val Loss: 0.7936, Val Acc: 0.9183

Learning Rate: 0.000004

```
{"model_id": "df0c9dfdf4c2409e86b73c5f4ef00fd1", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "dddbf2cf245944f29e6326d2e40dad77", "version_major": 2, "version_minor": 0}
```

Epoch 26/80 Summary:

Train Loss: 0.8475, Train Acc: 0.8939

Val Loss: 0.7835, Val Acc: 0.9222

Learning Rate: 0.000004

```
{"model_id": "70293ef685854e89b93f782c41c67b4a", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "df1ea72514184271a9611c0b2bc1614d", "version_major": 2, "version_minor": 0}
```


Epoch 27/80 Summary:

Train Loss: 0.8446, Train Acc: 0.8950

Val Loss: 0.7771, Val Acc: 0.9252

Learning Rate: 0.000004

```
{"model_id": "200e8130387d413badd9ef79d2ede817", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5b8747619e194e8aa431e1a352a065e4", "version_major": 2, "version_minor": 0}
```

Epoch 28/80 Summary:

Train Loss: 0.8385, Train Acc: 0.8982

Val Loss: 0.7754, Val Acc: 0.9294

Learning Rate: 0.000004

```
{"model_id": "707098f6f37a431c95b99af304772138", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "908f70e776f945d3bbe8c605860dd8ba", "version_major": 2, "version_minor": 0}
```

Epoch 29/80 Summary:

Train Loss: 0.8320, Train Acc: 0.9036

Val Loss: 0.7764, Val Acc: 0.9279

Learning Rate: 0.000004

```
{"model_id": "85d5035811bd425082a21d9d4d45246c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e8256a00bde548f7a15cd52b1cf4f706", "version_major": 2, "version_minor": 0}
```

Epoch 30/80 Summary:

Train Loss: 0.8326, Train Acc: 0.9026

Val Loss: 0.7805, Val Acc: 0.9252

Learning Rate: 0.000004

```
{"model_id": "c66aa0db3ee14e7aa0096e8d74f82433", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "16aa6ea4bfc14022b1a4e7bb85cc726a", "version_major": 2, "version_minor": 0}
```

Epoch 31/80 Summary:

Train Loss: 0.8255, Train Acc: 0.9026

Val Loss: 0.7639, Val Acc: 0.9303
Learning Rate: 0.000004

```
{"model_id": "a81d35e4adf14fd48a060c7842e4d22a", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "de119b736a784490b8b3289a75048b35", "version_major": 2, "version_minor": 0}
```

Epoch 32/80 Summary:

Train Loss: 0.8212, Train Acc: 0.9079
Val Loss: 0.7623, Val Acc: 0.9342
Learning Rate: 0.000004

```
{"model_id": "b552a195dd6345a58edb2da033b075a2", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "856f7b3464aa4d148ea77adf786fde9b", "version_major": 2, "version_minor": 0}
```

Epoch 33/80 Summary:

Train Loss: 0.8194, Train Acc: 0.9071
Val Loss: 0.7634, Val Acc: 0.9303
Learning Rate: 0.000004

```
{"model_id": "7f05e8c62dc042c28c875a2ba219b38d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "3fff37b812bb4830a3ffb0f7d60fd1d8", "version_major": 2, "version_minor": 0}
```

Epoch 34/80 Summary:

Train Loss: 0.8144, Train Acc: 0.9106
Val Loss: 0.7597, Val Acc: 0.9345
Learning Rate: 0.000004

```
{"model_id": "b8f8f49d0c4e4d9b9d201963dc79d9be", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b0a410a79de544d0820a36c17a810bc5", "version_major": 2, "version_minor": 0}
```

Epoch 35/80 Summary:

Train Loss: 0.8112, Train Acc: 0.9118
Val Loss: 0.7573, Val Acc: 0.9330
Learning Rate: 0.000004

```
{"model_id": "566474a53d2946c69e16304cbc41b2df", "version_major": 2, "version_minor": 0}
```

```
{"model_id":"eb60f3efaccd4d6593ebde86bf9c30f9","version_major":2,"version_minor":0}
```

Epoch 36/80 Summary:

Train Loss: 0.8098, Train Acc: 0.9139

Val Loss: 0.7585, Val Acc: 0.9366

Learning Rate: 0.000004

```
{"model_id":"9f28fb7cc5fb493195bff64ac2686d5c","version_major":2,"version_minor":0}
```

```
{"model_id":"fca4a45152ba4f668f7fa75f862864e0","version_major":2,"version_minor":0}
```

Epoch 37/80 Summary:

Train Loss: 0.8065, Train Acc: 0.9163

Val Loss: 0.7578, Val Acc: 0.9348

Learning Rate: 0.000004

```
{"model_id":"e61e9f1d106c4603aed5a71380bacc83","version_major":2,"version_minor":0}
```

```
{"model_id":"ba50470181ea41fa9632cad3ff0c84e0","version_major":2,"version_minor":0}
```

Epoch 38/80 Summary:

Train Loss: 0.8036, Train Acc: 0.9145

Val Loss: 0.7539, Val Acc: 0.9321

Learning Rate: 0.000004

```
{"model_id":"d059d49bfe3b45beb79d9855042b43da","version_major":2,"version_minor":0}
```

```
{"model_id":"cc08aedd34354d47afb2088b6f4ce9d1","version_major":2,"version_minor":0}
```

Epoch 39/80 Summary:

Train Loss: 0.8003, Train Acc: 0.9179

Val Loss: 0.7427, Val Acc: 0.9417

Learning Rate: 0.000004

```
{"model_id":"b248b441408649e9afb45e1fec579e68","version_major":2,"version_minor":0}
```

```
{"model_id":"f4bfb695f4064aala45f0a744f350e04","version_major":2,"version_minor":0}
```

Epoch 40/80 Summary:

Train Loss: 0.7918, Train Acc: 0.9197
Val Loss: 0.7438, Val Acc: 0.9405
Learning Rate: 0.000004

```
{"model_id": "052301b4a8b8471383f8f2006580715c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "8469ba6e73f74b8b900d2ff821e04f81", "version_major": 2, "version_minor": 0}
```

Epoch 41/80 Summary:
Train Loss: 0.7899, Train Acc: 0.9221
Val Loss: 0.7408, Val Acc: 0.9405
Learning Rate: 0.000004

```
{"model_id": "c9a5f75db1a948e580e49dbe4de2195c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e884047877c8482597539e87104c9643", "version_major": 2, "version_minor": 0}
```

Epoch 42/80 Summary:
Train Loss: 0.7924, Train Acc: 0.9206
Val Loss: 0.7481, Val Acc: 0.9381
Learning Rate: 0.000004

```
{"model_id": "b2bc8537f333483392e6deb22493b13d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "fd62c2cf0efd405ba800933ed49cd7d1", "version_major": 2, "version_minor": 0}
```

Epoch 43/80 Summary:
Train Loss: 0.7848, Train Acc: 0.9238
Val Loss: 0.7416, Val Acc: 0.9405
Learning Rate: 0.000004

```
{"model_id": "d08ef04136a4464b95b12e69c45a28c3", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c0c7db8ae05c43428bf163fcca248574", "version_major": 2, "version_minor": 0}
```

Epoch 44/80 Summary:
Train Loss: 0.7852, Train Acc: 0.9262
Val Loss: 0.7391, Val Acc: 0.9425
Learning Rate: 0.000004

```
{"model_id": "2f3f097263c14c53b7314299d4cad7e1", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e28caf6ca2d04dca9f1a667b4fe27824", "version_major": 2, "version_minor": 0}
```

Epoch 45/80 Summary:

Train Loss: 0.7795, Train Acc: 0.9263

Val Loss: 0.7306, Val Acc: 0.9449

Learning Rate: 0.000004

```
{"model_id": "75fc03b80c1b484c9e1a2079b89fdcca", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e20e5d31ea264e63b87976d6269ad10c", "version_major": 2, "version_minor": 0}
```

Epoch 46/80 Summary:

Train Loss: 0.7793, Train Acc: 0.9287

Val Loss: 0.7392, Val Acc: 0.9405

Learning Rate: 0.000004

```
{"model_id": "fc7cafd041ee498981b1e98361c09772", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0c77b8f7b89c4d9c8a64ee28dfaacc72e", "version_major": 2, "version_minor": 0}
```

Epoch 47/80 Summary:

Train Loss: 0.7778, Train Acc: 0.9292

Val Loss: 0.7347, Val Acc: 0.9446

Learning Rate: 0.000004

```
{"model_id": "8db375046bdc4b189d819097b66ea541", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0e5530168db1499d99262bc06ba147b7", "version_major": 2, "version_minor": 0}
```

Epoch 48/80 Summary:

Train Loss: 0.7775, Train Acc: 0.9301

Val Loss: 0.7284, Val Acc: 0.9476

Learning Rate: 0.000004

```
{"model_id": "388fce329eb34032869eb9f2896a3404", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "7463b1f8983c477787b25f72e1fbf2e3", "version_major": 2, "version_minor": 0}
```

Epoch 49/80 Summary:

Train Loss: 0.7720, Train Acc: 0.9318

Val Loss: 0.7315, Val Acc: 0.9431

Learning Rate: 0.000004

```
{"model_id": "59f7be67f6e64f7cad571af568b019cd", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "44e292ae464644369eae2616264b9e43", "version_major": 2, "version_minor": 0}
```

Epoch 50/80 Summary:

Train Loss: 0.7699, Train Acc: 0.9328

Val Loss: 0.7240, Val Acc: 0.9494

Learning Rate: 0.000004

```
{"model_id": "e5e8f13b17c24ebc9436687a7e40cfe1", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "35b190a86611468aa01bb3a34d7b986c", "version_major": 2, "version_minor": 0}
```

Epoch 51/80 Summary:

Train Loss: 0.7657, Train Acc: 0.9354

Val Loss: 0.7241, Val Acc: 0.9524

Learning Rate: 0.000004

```
{"model_id": "2b4b8938e1224ce88f051130fb0e0ec8", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "534f1350660142babdb388439ff5ae19", "version_major": 2, "version_minor": 0}
```

Epoch 52/80 Summary:

Train Loss: 0.7643, Train Acc: 0.9356

Val Loss: 0.7205, Val Acc: 0.9467

Learning Rate: 0.000004

```
{"model_id": "b46706f8ce4a43508dbd1e6e49436db5", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "977302b6ef1d4966be2510c83df1f865", "version_major": 2, "version_minor": 0}
```

Epoch 53/80 Summary:

Train Loss: 0.7658, Train Acc: 0.9329

Val Loss: 0.7216, Val Acc: 0.9515
Learning Rate: 0.000004

```
{"model_id": "e8418024c58b4c3cb80416f991413f52", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "7e005a68731b46dca93b9decda74743e", "version_major": 2, "version_minor": 0}
```

Epoch 54/80 Summary:
Train Loss: 0.7562, Train Acc: 0.9398
Val Loss: 0.7271, Val Acc: 0.9443
Learning Rate: 0.000004

```
{"model_id": "e11ebd55a25847eda26a36d879dec02b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6569e3fe246748aca5522a484b96b368", "version_major": 2, "version_minor": 0}
```

Epoch 55/80 Summary:
Train Loss: 0.7553, Train Acc: 0.9412
Val Loss: 0.7183, Val Acc: 0.9518
Learning Rate: 0.000004

```
{"model_id": "dcf9a43ef7c64229894944be5135eb73", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "39edabe8172f415183dbf3a967dc42c2", "version_major": 2, "version_minor": 0}
```

Epoch 56/80 Summary:
Train Loss: 0.7571, Train Acc: 0.9383
Val Loss: 0.7168, Val Acc: 0.9542
Learning Rate: 0.000004

```
{"model_id": "0eb974f4754747ff8cdc1dbe58b7c158", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "fde83b2748804801825220abb028823e", "version_major": 2, "version_minor": 0}
```

Epoch 57/80 Summary:
Train Loss: 0.7561, Train Acc: 0.9380
Val Loss: 0.7167, Val Acc: 0.9527
Learning Rate: 0.000004

```
{"model_id": "3b1b3c43dfe745158f06773eb055713f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "920f7fa776f341f484ddf900f21e5599", "version_major": 2, "version_minor": 0}
```

Epoch 58/80 Summary:

Train Loss: 0.7509, Train Acc: 0.9409

Val Loss: 0.7106, Val Acc: 0.9551

Learning Rate: 0.000004

```
{"model_id": "82c8e22fcaa34a28aa8db35b028eca87", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "77aa22d88f1f49a09cfade023fdd584b", "version_major": 2, "version_minor": 0}
```

Epoch 59/80 Summary:

Train Loss: 0.7511, Train Acc: 0.9421

Val Loss: 0.7133, Val Acc: 0.9506

Learning Rate: 0.000004

```
{"model_id": "45781552321e43cf97e6ea5e8d22e147", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "83b7f3005397445cade064a581ccf598", "version_major": 2, "version_minor": 0}
```

Epoch 60/80 Summary:

Train Loss: 0.7472, Train Acc: 0.9436

Val Loss: 0.7092, Val Acc: 0.9584

Learning Rate: 0.000004

```
{"model_id": "b3f3a9f1f6c94c89b09a8ffe29ea390c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "664a1749c5754c639eb70ff800135138", "version_major": 2, "version_minor": 0}
```

Epoch 61/80 Summary:

Train Loss: 0.7467, Train Acc: 0.9425

Val Loss: 0.7063, Val Acc: 0.9545

Learning Rate: 0.000004

```
{"model_id": "cd8cb7ee63974a519fd3d6c670ccbe88", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e18f812c56cc446d926e4e416aa69bc6", "version_major": 2, "version_minor": 0}
```

Epoch 62/80 Summary:

Train Loss: 0.7416, Train Acc: 0.9470
Val Loss: 0.7105, Val Acc: 0.9530
Learning Rate: 0.000004

```
{"model_id": "3d9b8a7888484daa97755b626f26938b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "843b7c69f7044da18f92e1195256059a", "version_major": 2, "version_minor": 0}
```

Epoch 63/80 Summary:
Train Loss: 0.7450, Train Acc: 0.9450
Val Loss: 0.7116, Val Acc: 0.9527
Learning Rate: 0.000004

```
{"model_id": "9b7c268136be4c6aa0b423c1b3f808ac", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f5806c2d6ba248ec9369a8db71239735", "version_major": 2, "version_minor": 0}
```

Epoch 64/80 Summary:
Train Loss: 0.7399, Train Acc: 0.9482
Val Loss: 0.7075, Val Acc: 0.9536
Learning Rate: 0.000004

```
{"model_id": "b9ca4ef473d44464be754ea0f3198b6a", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "ee1de921698842659674eef151f2d12b", "version_major": 2, "version_minor": 0}
```

Epoch 65/80 Summary:
Train Loss: 0.7384, Train Acc: 0.9487
Val Loss: 0.7080, Val Acc: 0.9539
Learning Rate: 0.000004

```
{"model_id": "c4b1cbc3061643aa8b7d455ba7996604", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "de271fb5d1124fa69ff3e23e14b4c48c", "version_major": 2, "version_minor": 0}
```

Epoch 66/80 Summary:
Train Loss: 0.7391, Train Acc: 0.9475
Val Loss: 0.7076, Val Acc: 0.9560
Learning Rate: 0.000004
Early stopping triggered!


Training completed in 2:11:16.915932


Plot training and validation accuracy

```
plt.figure(figsize=(12, 6))
plt.plot(history["train_acc"], label="Training Accuracy", marker="o")
plt.plot(history["val_acc"], label="Validation Accuracy", marker="o")
plt.title("Training and Validation Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.grid()
plt.show()
```

Plot training and validation loss

```
plt.figure(figsize=(12, 6))
plt.plot(history["train_loss"], label="Training Loss", marker="o")
plt.plot(history["val_loss"], label="Validation Loss", marker="o")
plt.title("Training and Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid()
plt.show()
```





```
def evaluate_model(model, data_loader, device):
    """
    Evaluates the accuracy of the model on the given DataLoader.

    Args:
        model: Trained PyTorch model.
        data_loader: DataLoader for the dataset to evaluate.
        device: Device to run the model on ('cpu' or 'cuda').

    Returns:
        accuracy: Accuracy of the model on the dataset.
    """
    model.eval() # Set the model to evaluation mode
    correct = 0
    total = 0

    with torch.no_grad(): # Disable gradient calculation for
        evaluation
        for inputs, labels in data_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            # Forward pass
            outputs = model(inputs)

            # Predictions
            predicted = outputs.argmax(dim=1)

            # Count correctly classified samples
            correct += (predicted == labels).sum().item()
```

```

        total += labels.size(0)

    # Calculate accuracy
    accuracy = correct / total
    return accuracy

# Evaluate the model
true_accuracy = evaluate_model(trained_model, test_loader, device)
print(f"Test Accuracy: {true_accuracy:.4f}")

Test Accuracy: 0.9545

def visualize_prediction(model, data, device, class_names):
    """
    Displays a test image along with its actual and predicted class
    probabilities.

    Args:
        model: Trained PyTorch model.
        data: A single (image, label) tuple from the dataset or
        DataLoader.
        device: The device to run the model on ('cpu' or 'cuda').
        class_names: List of class names corresponding to the labels.
    """
    model.eval() # Set the model to evaluation mode

    # Unpack the image and label
    img, label = data
    img, label = img.to(device), label.to(device)

    # Add batch dimension to the image if needed
    img_batch = img.unsqueeze(0)

    # Predict the class of the image
    with torch.no_grad():
        outputs = model(img_batch)
        probabilities = torch.nn.functional.softmax(outputs, dim=1)
        _, predicted = torch.max(outputs, 1)

    # Convert the image tensor to a NumPy array for visualization
    img_np = img.cpu().numpy().transpose((1, 2, 0)) # Convert to HWC
    format
    img_np = np.clip(img_np, 0, 1) # Ensure pixel values are in range
    [0, 1]

    # Get actual and predicted class names and probabilities
    actual_class = class_names[label.item()]
    predicted_class = class_names[predicted.item()]
    predicted_prob = probabilities[0, predicted.item()].item() * 100

```

```

# Plot the image with predicted and actual class
plt.figure(figsize=(6, 6))
plt.imshow(img_np)
plt.title(f"Actual: {actual_class} \nPredicted: {predicted_class}
({predicted_prob:.2f}%)", fontsize=14)
plt.axis('off')
plt.show()

# Visualizing a single image from the test_loader
for img, label in test_loader: # Get one batch of data
    visualize_prediction(trained_model, (img[0], label[0]), device,
CLASSES) # Visualize the first image
    break

```

Actual: Septoria_leaf_spot
Predicted: Septoria_leaf_spot (97.30%)



```

#Downloading the model's state_dict
torch.save(trained_model.state_dict(),

```

```
"/kaggle/working/trained_model.pth")  
  
#Also downloading the full model just in case  
torch.save(trained_model,  
"/kaggle/working/trained_model_complete.pth")
```